

ibaz5ckca

November 15, 2025

1 Task 2: Fully Convolutional Network (FCN)

by Dechathon Niamsa-ard [st126235]

Implement an FCN for semantic segmentation, train it on a small dataset, and analyze how architectural and training choices affect performance.

Instructions:

Dataset: - Use a subset of Pascal VOC, COCO, or a small custom dataset (10–20 images for quick runs). - Split into train/test. - Preprocess (resize, normalize, and convert masks to class indices).

Model Implementation: - Implement FCN-8s with skip connections from pool3 and pool4. - Use pretrained ResNet50 (remove final FC layers) as a backbone - Upsampling Method: 1) transpose convolution vs 2) bilinear interpolation.

Training - Loss: CrossEntropyLoss - Optimizer: Adam or SGD - Metrics: Mean IoU, pixel accuracy - Train for 20 epochs (or until convergence). - Log training curves.

Deliverable:

- Jupyter notebook with clear code and comments.
- Visualization of segmentation results (min. 3 test images).
- Table comparing transpose convolution vs bilinear interpolation.
- Summarize with visuals and short analysis.

1.1 Import Libraries and Setup

```
[1]: # Import necessary libraries
import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import Dataset, DataLoader
import torchvision
from torchvision import transforms
from torchvision.models import resnet50, ResNet50_Weights
import numpy as np
import matplotlib.pyplot as plt
from PIL import Image
import os
from tqdm import tqdm
```

```
import warnings
warnings.filterwarnings('ignore')

# Set random seed for reproducibility
torch.manual_seed(42)
np.random.seed(42)
```

```
[2]: # Set device to GPU if available
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print(f'Using device: {device}')
```

Using device: cuda

1.2 Dataset Preparation

Note: This notebook uses the local Pascal VOC dataset from the `datasets/task_2/` directory in the workspace.

Dataset Source: Pascal VOC 2012 Dataset

```
[3]: # Dataset path configuration
import os

workspace_path = os.getcwd()
print(f"Workspace directory: {workspace_path}")

# Try different possible dataset paths
possible_paths = [
    os.path.join(workspace_path, "datasets", "task_2", "VOC2012_train_val", "VOC2012_train_val")
]

path = None
for p in possible_paths:
    if os.path.exists(p) and os.path.exists(os.path.join(p, "JPEGImages")):
        path = p
        break

if path is None:
    raise FileNotFoundError("Pascal VOC dataset not found. Please ensure the dataset is in datasets/task_2/")

print(f"Dataset path: {path}")
print(f"Dataset found successfully")
```

Workspace directory: c:\Users\USER\Documents\Dechathon_N\AIT\Computer Vision\AIT-Computer-Vision-Assignment3
Dataset path: c:\Users\USER\Documents\Dechathon_N\AIT\Computer Vision\AIT-

```
Computer-Vision-Assignment3\datasets\task_2\VOC2012_train_val\VOC2012_train_val
Dataset found successfully
```

```
[4]: # Configure dataset directories
images_dir = os.path.join(path, 'JPEGImages')
masks_dir = os.path.join(path, 'SegmentationClass')

# Validate dataset structure
for dir_path, dir_name in [(images_dir, 'Images'), (masks_dir, 'Masks')]:
    if not os.path.exists(dir_path):
        available_dirs = os.listdir(path) if os.path.exists(path) else []
        raise FileNotFoundError(f"{dir_name} directory not found:\n{dir_path}\nAvailable: {available_dirs}")

print(f"Images: {images_dir}")
print(f"Masks: {masks_dir}")

# Count dataset files
image_files = [f for f in os.listdir(images_dir) if f.lower().endswith('.jpg')]
mask_files = [f for f in os.listdir(masks_dir) if f.lower().endswith('.png')]

print(f"Total images: {len(image_files)}")
print(f"Total masks: {len(mask_files)}")
print("Dataset validation complete")
```

```
Images: c:\Users\USER\Documents\Dechathon_N\AIT\Computer Vision\AIT-Computer-
Vision-
Assignment3\datasets\task_2\VOC2012_train_val\VOC2012_train_val\JPEGImages
Masks: c:\Users\USER\Documents\Dechathon_N\AIT\Computer Vision\AIT-Computer-
Vision-Assignment3\datasets\task_2\VOC2012_train_val\VOC2012_train_val\Segmentat-
ionClass
Total images: 17125
Total masks: 2913
Dataset validation complete
```

```
[5]: # Image selection and dataset splitting
import random
random.seed(42)

# Validate and collect valid image-mask pairs
all_images = [f for f in os.listdir(images_dir) if f.lower().endswith('.jpg')]

if not all_images:
    raise ValueError(f"No JPEG images found in {images_dir}")

valid_pairs = []
for img_file in all_images:
```

```

mask_file = img_file.replace('.jpg', '.png')
if os.path.exists(os.path.join(masks_dir, mask_file)):
    valid_pairs.append(img_file)

print(f"Found {len(all_images)} images, {len(valid_pairs)} with corresponding masks")

if len(valid_pairs) < 10:
    raise ValueError(f"Insufficient valid pairs ({len(valid_pairs)}). Need at least 10 for training.")

# Select training subset (up to 100 images)
max_samples = min(1000, len(valid_pairs))
selected_images = random.sample(valid_pairs, max_samples)

# Train/test split (80/20)
split_idx = max(1, int(0.8 * len(selected_images)))
train_images = selected_images[:split_idx]
test_images = selected_images[split_idx:]

print("Dataset split:")
print(f" Total selected: {len(selected_images)}")
print(f" Training: {len(train_images)}")
print(f" Testing: {len(test_images)}")
print(f" Sample files: {selected_images[:3]}")

```

```

Found 17125 images, 2913 with corresponding masks
Dataset split:
Total selected: 1000
Training: 800
Testing: 200
Sample files: ['2011_000658.jpg', '2007_007355.jpg', '2007_001698.jpg']

```

[6]:

```

# Pascal VOC color map for visualization
VOC_COLORMAP = [[0, 0, 0], [128, 0, 0], [0, 128, 0], [128, 128, 0],
                 [0, 0, 128], [128, 0, 128], [0, 128, 128], [128, 128, 128],
                 [64, 0, 0], [192, 0, 0], [64, 128, 0], [192, 128, 0],
                 [64, 0, 128], [192, 0, 128], [64, 128, 128], [192, 128, 128],
                 [0, 64, 0], [128, 64, 0], [0, 192, 0], [128, 192, 0],
                 [0, 64, 128]]

VOC_CLASSES = ['background', 'aeroplane', 'bicycle', 'bird', 'boat',
               'bottle', 'bus', 'car', 'cat', 'chair', 'cow',
               'diningtable', 'dog', 'horse', 'motorbike', 'person',
               'potted plant', 'sheep', 'sofa', 'train', 'tv/monitor']

# Create colormap dictionary

```

```

colormap2label = np.zeros(256 ** 3, dtype=np.uint8)
for i, colormap in enumerate(VOC_COLORMAP):
    colormap2label[(colormap[0] * 256 + colormap[1]) * 256 + colormap[2]] = i

def voc_label_indices(colormap):
    """Convert RGB mask to class indices"""
    colormap = np.array(colormap).astype(np.int32)
    idx = ((colormap[:, :, 0] * 256 + colormap[:, :, 1]) * 256 + colormap[:, :, 2])
    return colormap2label[idx]

```

```

[7]: class VOCSegmentationDataset(Dataset):
    """Pascal VOC Dataset for semantic segmentation."""

    def __init__(self, images_dir, masks_dir, image_list, target_size=(256, 256)):
        self.images_dir = images_dir
        self.masks_dir = masks_dir
        self.image_list = image_list
        self.target_size = target_size

        # ImageNet normalization parameters
        self.normalize = transforms.Normalize(
            mean=[0.485, 0.456, 0.406],
            std=[0.229, 0.224, 0.225]
        )

        # Validate all files exist
        self._validate_files()

    def _validate_files(self):
        """Validate that all image-mask pairs exist."""
        missing_files = []
        for img_name in self.image_list:
            img_path = os.path.join(self.images_dir, img_name)
            mask_path = os.path.join(self.masks_dir, img_name.replace('.jpg', '.png'))

            if not os.path.exists(img_path):
                missing_files.append(f"Image: {img_path}")
            if not os.path.exists(mask_path):
                missing_files.append(f"Mask: {mask_path}")

        if missing_files:
            raise FileNotFoundError(f"Missing files: {missing_files[:5]}...")

    def __len__(self):

```

```

        return len(self.image_list)

    def __getitem__(self, idx):
        img_name = self.image_list[idx]
        img_path = os.path.join(self.images_dir, img_name)
        mask_path = os.path.join(self.masks_dir, img_name.replace('.jpg', '.  
png'))

    try:
        # Load and preprocess image
        image = Image.open(img_path).convert('RGB')
        mask = Image.open(mask_path).convert('RGB')

        # Resize with appropriate interpolation
        image = image.resize(self.target_size, Image.BILINEAR)
        mask = mask.resize(self.target_size, Image.NEAREST)

        # Convert to arrays
        image_array = np.array(image, dtype=np.float32)
        mask_array = np.array(mask)

        # Convert mask to class indices
        mask_indices = voc_label_indices(mask_array)

        # Convert to tensors
        image_tensor = torch.from_numpy(image_array).permute(2, 0, 1) / 255.  
0
        mask_tensor = torch.from_numpy(mask_indices).long()

        # Normalize image
        image_tensor = self.normalize(image_tensor)

        return image_tensor, mask_tensor

    except Exception as e:
        raise RuntimeError(f"Error loading {img_name}: {e}")

print("VOC Dataset class ready")

```

VOC Dataset class ready

```
[8]: # Create datasets and data loaders
train_dataset = VOCSegmentationDataset(images_dir, masks_dir, train_images)
test_dataset = VOCSegmentationDataset(images_dir, masks_dir, test_images)

train_loader = DataLoader(  
    train_dataset,
```

```

batch_size=2,
shuffle=True,
num_workers=0, # Set to 0 for Windows compatibility
pin_memory=True if torch.cuda.is_available() else False
)
test_loader = DataLoader(
    test_dataset,
    batch_size=2,
    shuffle=False,
    num_workers=0,
    pin_memory=True if torch.cuda.is_available() else False
)

print(f"Dataset loaded:")
print(f" Training: {len(train_dataset)} samples")
print(f" Testing: {len(test_dataset)} samples")

# Validate dataset loading
try:
    sample_img, sample_mask = train_dataset[0]
    print(f"Sample validation:")
    print(f" Image shape: {sample_img.shape}")
    print(f" Mask shape: {sample_mask.shape}")
    print(f" Classes in sample: {torch.unique(sample_mask).tolist()}")
except Exception as e:
    print(f"Dataset loading error: {e}")

```

```

Dataset loaded:
Training: 800 samples
Testing: 200 samples
Sample validation:
Image shape: torch.Size([3, 256, 256])
Mask shape: torch.Size([256, 256])
Classes in sample: [0, 7]
Classes in sample: [0, 7]

```

1.3 Model Implementation

```

[9]: class FCN8s(nn.Module):
    """
    Fully Convolutional Network (8s) for semantic segmentation.
    Uses ResNet50 as backbone with skip connections from pool3 and pool4.
    Supports both transpose convolution and bilinear interpolation for
    upsampling.
    """

    def __init__(self, num_classes=21, upsampling='transpose'):

```

```

super(FCN8s, self).__init__()
self.num_classes = num_classes
self.upsampling = upsampling

# Load pretrained ResNet50 and extract feature layers
resnet = resnet50(weights=ResNet50_Weights.DEFAULT)

# Extract different levels of features for skip connections
# pool3: after layer2 (512 channels, 1/8 resolution)
# pool4: after layer3 (1024 channels, 1/16 resolution)
# pool5: after layer4 (2048 channels, 1/32 resolution)

self.layer0 = nn.Sequential(resnet.conv1, resnet.bn1, resnet.relu, ▾
                           resnet.maxpool)
self.layer1 = resnet.layer1 # 1/4 resolution, 256 channels
self.layer2 = resnet.layer2 # 1/8 resolution, 512 channels (pool3)
self.layer3 = resnet.layer3 # 1/16 resolution, 1024 channels (pool4)
self.layer4 = resnet.layer4 # 1/32 resolution, 2048 channels (pool5)

# Classifier heads for different scales
self.classifier_pool5 = nn.Conv2d(2048, num_classes, kernel_size=1) # ▾
↪1/32
self.classifier_pool4 = nn.Conv2d(1024, num_classes, kernel_size=1) # ▾
↪1/16
self.classifier_pool3 = nn.Conv2d(512, num_classes, kernel_size=1) # ▾
↪1/8

# Upsampling layers for FCN-8s
if upsampling == 'transpose':
    # 2x upsampling from pool5 to pool4 resolution
    self.upsample_pool5 = nn.ConvTranspose2d(
        num_classes, num_classes, kernel_size=4, stride=2, padding=1, ▾
        bias=False
    )
    # 2x upsampling from fused pool4 to pool3 resolution
    self.upsample_pool4 = nn.ConvTranspose2d(
        num_classes, num_classes, kernel_size=4, stride=2, padding=1, ▾
        bias=False
    )
    # 8x upsampling from fused pool3 to original resolution
    self.upsample_final = nn.ConvTranspose2d(
        num_classes, num_classes, kernel_size=16, stride=8, padding=4, ▾
        bias=False
    )

# Initialize with bilinear weights

```

```

        self._init_transpose_weights()
    else: # bilinear interpolation
        self.upsample_2x = nn.Upsample(scale_factor=2, mode='bilinear',  

        ↪align_corners=False)
        self.upsample_8x = nn.Upsample(scale_factor=8, mode='bilinear',  

        ↪align_corners=False)

    def _init_transpose_weights(self):
        """Initialize transpose convolution layers with bilinear weights"""
        def init_layer(layer):
            kernel_size = layer.kernel_size[0]
            stride = layer.stride[0]

            # Create bilinear kernel
            factor = (kernel_size + 1) // 2
            if kernel_size % 2 == 1:
                center = factor - 1
            else:
                center = factor - 0.5

            og = np.ogrid[:kernel_size, :kernel_size]
            filt = (1 - abs(og[0] - center) / factor) * (1 - abs(og[1] -  

            ↪center) / factor)
            weight = np.zeros((self.num_classes, self.num_classes, kernel_size,  

            ↪kernel_size))

            for i in range(self.num_classes):
                weight[i, i, :, :] = filt

            layer.weight.data.copy_(torch.from_numpy(weight))

        # Initialize all transpose convolution layers
        init_layer(self.upsample_pool5)
        init_layer(self.upsample_pool4)
        init_layer(self.upsample_final)

    def forward(self, x):
        # Save input size for final resize
        input_size = x.shape[2:]

        # Extract multi-scale features
        x = self.layer0(x)      # 1/4 resolution
        x = self.layer1(x)      # 1/4 resolution, 256 channels
        pool3 = self.layer2(x)  # 1/8 resolution, 512 channels
        pool4 = self.layer3(pool3) # 1/16 resolution, 1024 channels
        pool5 = self.layer4(pool4) # 1/32 resolution, 2048 channels

```

```

# Apply classifiers to each scale
score_pool5 = self.classifier_pool5(pool5) # 1/32 resolution
score_pool4 = self.classifier_pool4(pool4) # 1/16 resolution
score_pool3 = self.classifier_pool3(pool3) # 1/8 resolution

if self.upsampling == 'transpose':
    # FCN-8s with transpose convolutions
    # Upsample pool5 by 2x and fuse with pool4
    upscore_pool5 = self.upsample_pool5(score_pool5)
    # Ensure size compatibility for addition
    if upscore_pool5.shape[2:] != score_pool4.shape[2:]:
        upscore_pool5 = nn.functional.interpolate(
            upscore_pool5, size=score_pool4.shape[2:], mode='bilinear', □
    ↪align_corners=False
)
fuse_pool4 = upscore_pool5 + score_pool4

# Upsample fused pool4 by 2x and fuse with pool3
upscore_pool4 = self.upsample_pool4(fuse_pool4)
if upscore_pool4.shape[2:] != score_pool3.shape[2:]:
    upscore_pool4 = nn.functional.interpolate(
        upscore_pool4, size=score_pool3.shape[2:], mode='bilinear', □
    ↪align_corners=False
)
fuse_pool3 = upscore_pool4 + score_pool3

# Final 8x upsampling to original resolution
output = self.upsample_final(fuse_pool3)

else: # bilinear interpolation
    # FCN-8s with bilinear interpolation
    # Upsample pool5 by 2x and fuse with pool4
    upscore_pool5 = self.upsample_2x(score_pool5)
    if upscore_pool5.shape[2:] != score_pool4.shape[2:]:
        upscore_pool5 = nn.functional.interpolate(
            upscore_pool5, size=score_pool4.shape[2:], mode='bilinear', □
    ↪align_corners=False
)
    fuse_pool4 = upscore_pool5 + score_pool4

    # Upsample fused pool4 by 2x and fuse with pool3
    upscore_pool4 = self.upsample_2x(fuse_pool4)
    if upscore_pool4.shape[2:] != score_pool3.shape[2:]:
        upscore_pool4 = nn.functional.interpolate(
            upscore_pool4, size=score_pool3.shape[2:], mode='bilinear', □
    ↪align_corners=False
)

```

```

        fuse_pool3 = upscore_pool4 + score_pool3

        # Final 8x upsampling to original resolution
        output = self.upsample_8x(fuse_pool3)

        # Ensure exact size match with input
        if output.shape[2:] != input_size:
            output = nn.functional.interpolate(output, size=input_size, mode='bilinear', align_corners=False)

    return output

print("FCN8s model class implemented successfully")

```

FCN8s model class implemented successfully

1.4 Training and Evaluation Functions

```
[10]: def calculate_iou(pred, target, num_classes=21):
    """Calculate mean Intersection over Union (IoU)"""
    ious = []
    pred = pred.view(-1)
    target = target.view(-1)

    for cls in range(num_classes):
        pred_cls = pred == cls
        target_cls = target == cls

        if target_cls.long().sum().item() == 0:
            # If class is not present in ground truth, skip
            continue

        intersection = (pred_cls & target_cls).long().sum().item()
        union = (pred_cls | target_cls).long().sum().item()

        if union == 0:
            iou = 1.0 # Perfect score when both are empty
        else:
            iou = intersection / union
        ious.append(iou)

    return np.mean(ious) if ious else 0.0

def calculate_pixel_accuracy(pred, target):
    """Calculate pixel accuracy"""
    correct = (pred == target).sum().item()
    total = target.numel()
```

```

    return correct / total

def train_epoch(model, train_loader, criterion, optimizer, device):
    """Train model for one epoch"""
    model.train()
    total_loss = 0.0
    total_iou = 0.0
    total_accuracy = 0.0
    num_batches = 0

    progress_bar = tqdm(train_loader, desc="Training")

    for images, masks in progress_bar:
        images, masks = images.to(device), masks.to(device)

        optimizer.zero_grad()

        # Forward pass
        outputs = model(images)
        loss = criterion(outputs, masks)

        # Backward pass
        loss.backward()
        optimizer.step()

        # Calculate metrics
        _, predicted = torch.max(outputs, 1)
        batch_iou = calculate_iou(predicted.cpu(), masks.cpu())
        batch_accuracy = calculate_pixel_accuracy(predicted.cpu(), masks.cpu())

        total_loss += loss.item()
        total_iou += batch_iou
        total_accuracy += batch_accuracy
        num_batches += 1

        # Update progress bar
        progress_bar.set_postfix({
            'Loss': f'{loss.item():.4f}',
            'IoU': f'{batch_iou:.4f}',
            'Acc': f'{batch_accuracy:.4f}'
        })

    return total_loss / num_batches, total_iou / num_batches, total_accuracy / num_batches

def evaluate_model(model, test_loader, criterion, device):
    """Evaluate model on test set"""

```

```

model.eval()
total_loss = 0.0
total_iou = 0.0
total_accuracy = 0.0
num_batches = 0

with torch.no_grad():
    progress_bar = tqdm(test_loader, desc="Evaluating")

    for images, masks in progress_bar:
        images, masks = images.to(device), masks.to(device)

        # Forward pass
        outputs = model(images)
        loss = criterion(outputs, masks)

        # Calculate metrics
        _, predicted = torch.max(outputs, 1)
        batch_iou = calculate_iou(predicted.cpu(), masks.cpu())
        batch_accuracy = calculate_pixel_accuracy(predicted.cpu(), masks.
cpu())

        total_loss += loss.item()
        total_iou += batch_iou
        total_accuracy += batch_accuracy
        num_batches += 1

        # Update progress bar
        progress_bar.set_postfix({
            'Loss': f'{loss.item():.4f}',
            'IoU': f'{batch_iou:.4f}',
            'Acc': f'{batch_accuracy:.4f}'
        })

    return total_loss / num_batches, total_iou / num_batches, total_accuracy / num_batches

```



```

def train_model(model, train_loader, test_loader, num_epochs=20, learning_rate=1e-4):
    """Complete training loop"""
    criterion = nn.CrossEntropyLoss(ignore_index=255) # Ignore void pixels
    optimizer = optim.Adam(model.parameters(), lr=learning_rate, weight_decay=1e-4)

    # Learning rate scheduler
    scheduler = optim.lr_scheduler.StepLR(optimizer, step_size=10, gamma=0.1)

```

```

# History tracking
history = {
    'train_loss': [], 'train_iou': [], 'train_accuracy': [],
    'val_loss': [], 'val_iou': [], 'val_accuracy': []
}

best_iou = 0.0

print(f"Starting training for {num_epochs} epochs...")
print(f"Model: {model.__class__.__name__} with {model.upsampling} upsampling")
print("-" * 60)

for epoch in range(num_epochs):
    print(f"\nEpoch {epoch+1}/{num_epochs}")

    # Training phase
    train_loss, train_iou, train_acc = train_epoch(model, train_loader,
criterion, optimizer, device)

    # Validation phase
    val_loss, val_iou, val_acc = evaluate_model(model, test_loader,
criterion, device)

    # Update learning rate
    scheduler.step()

    # Store history
    history['train_loss'].append(train_loss)
    history['train_iou'].append(train_iou)
    history['train_accuracy'].append(train_acc)
    history['val_loss'].append(val_loss)
    history['val_iou'].append(val_iou)
    history['val_accuracy'].append(val_acc)

    # Save best model
    if val_iou > best_iou:
        best_iou = val_iou
        torch.save(model.state_dict(), f'models/task_2/fcn8s_{model.
upsampling}_best.pth')

    # Print epoch summary
    print(f"Train - Loss: {train_loss:.4f}, IoU: {train_iou:.4f}, Acc: {train_acc:.4f}")
    print(f"Val - Loss: {val_loss:.4f}, IoU: {val_iou:.4f}, Acc: {val_acc:.4f}")

```

```

    print(f"Best Val IoU: {best_iou:.4f}")

    # Save final model
    torch.save(model.state_dict(), f'models/task_2/fcn8s_{model.
    ↪upsampling}_final.pth')

    print(f"\nTraining completed! Best validation IoU: {best_iou:.4f}")
    return history

print("Training and evaluation functions implemented successfully")

```

Training and evaluation functions implemented successfully

1.5 Train Models with Different Upsampling Methods

1.5.1 1. FCN-8s with Transpose Convolution

```
[11]: # Create and train FCN with transpose convolution
print("=" * 80)
print("TRAINING FCN-8s WITH TRANSPOSE CONVOLUTION")
print("=" * 80)

# Initialize model
fcn_transpose = FCN8s(num_classes=21, upsampling='transpose').to(device)

# Check if pretrained model exists
transpose_model_path = 'models/task_2/fcn8s_transpose_best.pth'
if os.path.exists(transpose_model_path):
    print("Loading pretrained transpose convolution model...")
    fcn_transpose.load_state_dict(torch.load(transpose_model_path, ↪
    ↪map_location=device))

    # Quick evaluation
    val_loss, val_iou, val_acc = evaluate_model(fcn_transpose, test_loader, nn.
    ↪CrossEntropyLoss(ignore_index=255), device)
    print(f"Loaded model performance - Loss: {val_loss:.4f}, IoU: {val_iou:.
    ↪4f}, Acc: {val_acc:.4f}")

    # Create dummy history for consistency
    transpose_history = {
        'train_loss': [], 'train_iou': [], 'train_accuracy': [],
        'val_loss': [], 'val_iou': [], 'val_accuracy': []
    }
else:
    print("Training new transpose convolution model...")
    transpose_history = train_model(fcn_transpose, train_loader, test_loader, ↪
    ↪num_epochs=30, learning_rate=1e-4)
```

```
print("\nTranspose convolution training completed!")
```

```
=====
TRAINING FCN-8s WITH TRANSPOSE CONVOLUTION
=====
```

```
Training new transpose convolution model...
Starting training for 30 epochs...
Model: FCN8s with transpose upsampling
```

```
-----
Epoch 1/30
```

```
Training new transpose convolution model...
Starting training for 30 epochs...
Model: FCN8s with transpose upsampling
```

```
-----
Epoch 1/30
```

```
Training: 100%|      | 400/400 [00:16<00:00, 24.38it/s, Loss=1.0036,
IoU=0.3168, Acc=0.7755]
Training: 100%|      | 400/400 [00:16<00:00, 24.38it/s, Loss=1.0036,
IoU=0.3168, Acc=0.7755]
Evaluating: 100%|     | 100/100 [00:01<00:00, 51.79it/s, Loss=0.4711,
IoU=0.4386, Acc=0.8512]
```

```
Train - Loss: 1.0048, IoU: 0.3181, Acc: 0.7669
Val   - Loss: 0.6211, IoU: 0.4447, Acc: 0.8322
Best Val IoU: 0.4447
```

```
-----
Epoch 2/30
```

```
Training: 100%|      | 400/400 [00:14<00:00, 27.30it/s, Loss=0.2680,
IoU=0.8125, Acc=0.9291]
Training: 100%|      | 400/400 [00:14<00:00, 27.30it/s, Loss=0.2680,
IoU=0.8125, Acc=0.9291]
Evaluating: 100%|     | 100/100 [00:01<00:00, 52.85it/s, Loss=0.2210,
IoU=0.6057, Acc=0.9332]
```

```
Train - Loss: 0.6432, IoU: 0.4687, Acc: 0.8317
Val   - Loss: 0.5239, IoU: 0.5174, Acc: 0.8517
Best Val IoU: 0.5174
```

```
-----
Epoch 3/30
```

```
Training: 100%|      | 400/400 [00:14<00:00, 27.25it/s, Loss=0.3522,
IoU=0.5817, Acc=0.9083]
Training: 100%|      | 400/400 [00:14<00:00, 27.25it/s, Loss=0.3522,
```

```
[IoU=0.5817, Acc=0.9083]
Evaluating: 100%|      | 100/100 [00:01<00:00, 50.93it/s, Loss=0.1671,
IoU=0.6541, Acc=0.9528]
```

```
Train - Loss: 0.4775, IoU: 0.5623, Acc: 0.8700
Val   - Loss: 0.4788, IoU: 0.5589, Acc: 0.8655
Best Val IoU: 0.5589
```

Epoch 4/30

```
Training: 100%|      | 400/400 [00:14<00:00, 27.36it/s, Loss=1.7489,
IoU=0.4804, Acc=0.5817]
Training: 100%|      | 400/400 [00:14<00:00, 27.36it/s, Loss=1.7489,
IoU=0.4804, Acc=0.5817]
Evaluating: 100%|      | 100/100 [00:01<00:00, 52.23it/s, Loss=0.1510,
IoU=0.6480, Acc=0.9510]
```

```
Train - Loss: 0.3537, IoU: 0.6323, Acc: 0.9005
Val   - Loss: 0.4581, IoU: 0.5699, Acc: 0.8718
Best Val IoU: 0.5699
```

Epoch 5/30

```
Training: 100%|      | 400/400 [00:14<00:00, 27.31it/s, Loss=0.2771,
IoU=0.4571, Acc=0.9186]
Training: 100%|      | 400/400 [00:14<00:00, 27.31it/s, Loss=0.2771,
IoU=0.4571, Acc=0.9186]
Evaluating: 100%|      | 100/100 [00:01<00:00, 53.47it/s, Loss=0.1246,
IoU=0.6651, Acc=0.9584]
```

```
Train - Loss: 0.2963, IoU: 0.6652, Acc: 0.9150
Val   - Loss: 0.4385, IoU: 0.5781, Acc: 0.8728
Best Val IoU: 0.5781
```

Epoch 6/30

```
Training: 100%|      | 400/400 [00:14<00:00, 27.47it/s, Loss=0.4525,
IoU=0.7164, Acc=0.8633]
Training: 100%|      | 400/400 [00:14<00:00, 27.47it/s, Loss=0.4525,
IoU=0.7164, Acc=0.8633]
Evaluating: 100%|      | 100/100 [00:01<00:00, 52.93it/s, Loss=0.1238,
IoU=0.6645, Acc=0.9586]
Evaluating: 100%|      | 100/100 [00:01<00:00, 52.93it/s, Loss=0.1238,
IoU=0.6645, Acc=0.9586]
```

```
Train - Loss: 0.2260, IoU: 0.7011, Acc: 0.9319
Val   - Loss: 0.4302, IoU: 0.5711, Acc: 0.8745
Best Val IoU: 0.5781
```

Epoch 7/30

```
Training: 100%|      | 400/400 [00:14<00:00, 27.41it/s, Loss=0.1576,
IoU=0.6618, Acc=0.9455]
Training: 100%|      | 400/400 [00:14<00:00, 27.41it/s, Loss=0.1576,
IoU=0.6618, Acc=0.9455]
Evaluating: 100%|     | 100/100 [00:01<00:00, 53.05it/s, Loss=0.1228,
IoU=0.6629, Acc=0.9573]
Evaluating: 100%|     | 100/100 [00:01<00:00, 53.05it/s, Loss=0.1228,
IoU=0.6629, Acc=0.9573]

Train - Loss: 0.2012, IoU: 0.7250, Acc: 0.9390
Val   - Loss: 0.4464, IoU: 0.5873, Acc: 0.8734
Best Val IoU: 0.5873
```

Epoch 8/30

```
Training: 100%|      | 400/400 [00:14<00:00, 27.32it/s, Loss=0.2230,
IoU=0.6840, Acc=0.9291]
Training: 100%|      | 400/400 [00:14<00:00, 27.32it/s, Loss=0.2230,
IoU=0.6840, Acc=0.9291]
Evaluating: 100%|     | 100/100 [00:01<00:00, 53.50it/s, Loss=0.1576,
IoU=0.6463, Acc=0.9463]
Evaluating: 100%|     | 100/100 [00:01<00:00, 53.50it/s, Loss=0.1576,
IoU=0.6463, Acc=0.9463]

Train - Loss: 0.1747, IoU: 0.7398, Acc: 0.9461
Val   - Loss: 0.4509, IoU: 0.5886, Acc: 0.8704
Best Val IoU: 0.5886
```

Epoch 9/30

```
Training: 100%|      | 400/400 [00:14<00:00, 27.48it/s, Loss=0.0655,
IoU=0.7823, Acc=0.9828]
Training: 100%|      | 400/400 [00:14<00:00, 27.48it/s, Loss=0.0655,
IoU=0.7823, Acc=0.9828]
Evaluating: 100%|     | 100/100 [00:01<00:00, 52.43it/s, Loss=0.1310,
IoU=0.6549, Acc=0.9529]
```

```
Train - Loss: 0.1528, IoU: 0.7544, Acc: 0.9512
Val   - Loss: 0.4387, IoU: 0.5951, Acc: 0.8784
Best Val IoU: 0.5951
```

Epoch 10/30

```
Training: 100%|      | 400/400 [00:14<00:00, 27.38it/s, Loss=0.2251,
IoU=0.8699, Acc=0.9328]
Training: 100%|      | 400/400 [00:14<00:00, 27.38it/s, Loss=0.2251,
IoU=0.8699, Acc=0.9328]
```

```
Evaluating: 100%|      | 100/100 [00:01<00:00, 53.65it/s, Loss=0.1456,
IoU=0.6488, Acc=0.9501]
Evaluating: 100%|      | 100/100 [00:01<00:00, 53.65it/s, Loss=0.1456,
IoU=0.6488, Acc=0.9501]
Train - Loss: 0.1541, IoU: 0.7562, Acc: 0.9507
Val   - Loss: 0.4485, IoU: 0.5923, Acc: 0.8696
Best Val IoU: 0.5951
```

Epoch 11/30

```
Training: 100%|      | 400/400 [00:14<00:00, 27.29it/s, Loss=0.0382,
IoU=0.9104, Acc=0.9889]
Training: 100%|      | 400/400 [00:14<00:00, 27.29it/s, Loss=0.0382,
IoU=0.9104, Acc=0.9889]
Evaluating: 100%|      | 100/100 [00:01<00:00, 51.93it/s, Loss=0.1573,
IoU=0.6235, Acc=0.9377]
Evaluating: 100%|      | 100/100 [00:01<00:00, 51.93it/s, Loss=0.1573,
IoU=0.6235, Acc=0.9377]
Train - Loss: 0.1312, IoU: 0.7682, Acc: 0.9579
Val   - Loss: 0.4436, IoU: 0.5779, Acc: 0.8751
Best Val IoU: 0.5951
```

Epoch 12/30

```
Training: 100%|      | 400/400 [00:14<00:00, 27.48it/s, Loss=0.1027,
IoU=0.6662, Acc=0.9618]
Training: 100%|      | 400/400 [00:14<00:00, 27.48it/s, Loss=0.1027,
IoU=0.6662, Acc=0.9618]
Evaluating: 100%|      | 100/100 [00:01<00:00, 51.96it/s, Loss=0.1487,
IoU=0.6284, Acc=0.9412]
Evaluating: 100%|      | 100/100 [00:01<00:00, 51.96it/s, Loss=0.1487,
IoU=0.6284, Acc=0.9412]
Train - Loss: 0.1183, IoU: 0.7757, Acc: 0.9603
Val   - Loss: 0.4500, IoU: 0.5714, Acc: 0.8732
Best Val IoU: 0.5951
```

Epoch 13/30

```
Training: 100%|      | 400/400 [00:14<00:00, 27.65it/s, Loss=0.0647,
IoU=0.7586, Acc=0.9775]
Training: 100%|      | 400/400 [00:14<00:00, 27.65it/s, Loss=0.0647,
IoU=0.7586, Acc=0.9775]
Evaluating: 100%|      | 100/100 [00:01<00:00, 54.26it/s, Loss=0.1433,
IoU=0.6367, Acc=0.9454]
Evaluating: 100%|      | 100/100 [00:01<00:00, 54.26it/s, Loss=0.1433,
IoU=0.6367, Acc=0.9454]
Train - Loss: 0.1154, IoU: 0.7877, Acc: 0.9609
```

Val - Loss: 0.4388, IoU: 0.5732, Acc: 0.8746
Best Val IoU: 0.5951

Epoch 14/30

Training: 100% | 400/400 [00:14<00:00, 27.53it/s, Loss=0.0385,
IoU=0.9367, Acc=0.9890]
Training: 100% | 400/400 [00:14<00:00, 27.53it/s, Loss=0.0385,
IoU=0.9367, Acc=0.9890]
Evaluating: 100% | 100/100 [00:01<00:00, 53.22it/s, Loss=0.1617,
IoU=0.6197, Acc=0.9364]
Evaluating: 100% | 100/100 [00:01<00:00, 53.22it/s, Loss=0.1617,
IoU=0.6197, Acc=0.9364]

Train - Loss: 0.1072, IoU: 0.7892, Acc: 0.9637
Val - Loss: 0.4383, IoU: 0.5823, Acc: 0.8773
Best Val IoU: 0.5951

Epoch 15/30

Training: 100% | 400/400 [00:14<00:00, 27.61it/s, Loss=0.0540,
IoU=0.9318, Acc=0.9796]
Training: 100% | 400/400 [00:14<00:00, 27.61it/s, Loss=0.0540,
IoU=0.9318, Acc=0.9796]
Evaluating: 100% | 100/100 [00:01<00:00, 53.35it/s, Loss=0.1464,
IoU=0.6290, Acc=0.9413]
Evaluating: 100% | 100/100 [00:01<00:00, 53.35it/s, Loss=0.1464,
IoU=0.6290, Acc=0.9413]

Train - Loss: 0.1050, IoU: 0.7946, Acc: 0.9638
Val - Loss: 0.4386, IoU: 0.5800, Acc: 0.8769
Best Val IoU: 0.5951

Epoch 16/30

Training: 100% | 400/400 [00:14<00:00, 27.62it/s, Loss=0.0966,
IoU=0.9182, Acc=0.9687]
Training: 100% | 400/400 [00:14<00:00, 27.62it/s, Loss=0.0966,
IoU=0.9182, Acc=0.9687]
Evaluating: 100% | 100/100 [00:01<00:00, 52.50it/s, Loss=0.1345,
IoU=0.6437, Acc=0.9486]
Evaluating: 100% | 100/100 [00:01<00:00, 52.50it/s, Loss=0.1345,
IoU=0.6437, Acc=0.9486]

Train - Loss: 0.0993, IoU: 0.7955, Acc: 0.9654
Val - Loss: 0.4378, IoU: 0.5845, Acc: 0.8798
Best Val IoU: 0.5951

Epoch 17/30

Training: 100% | 400/400 [00:14<00:00, 27.43it/s, Loss=0.1911,

```
[IoU=0.7858, Acc=0.9344]
Training: 100%|      | 400/400 [00:14<00:00, 27.43it/s, Loss=0.1911,
IoU=0.7858, Acc=0.9344]
Evaluating: 100%|      | 100/100 [00:01<00:00, 52.54it/s, Loss=0.1465,
IoU=0.6283, Acc=0.9418]
Evaluating: 100%|      | 100/100 [00:01<00:00, 52.54it/s, Loss=0.1465,
IoU=0.6283, Acc=0.9418]

Train - Loss: 0.0991, IoU: 0.7983, Acc: 0.9653
Val   - Loss: 0.4463, IoU: 0.5800, Acc: 0.8779
Best Val IoU: 0.5951
```

Epoch 18/30

```
Training: 100%|      | 400/400 [00:14<00:00, 27.35it/s, Loss=0.1104,
IoU=0.7007, Acc=0.9619]
Training: 100%|      | 400/400 [00:14<00:00, 27.35it/s, Loss=0.1104,
IoU=0.7007, Acc=0.9619]
Evaluating: 100%|      | 100/100 [00:02<00:00, 48.13it/s, Loss=0.1356,
IoU=0.6406, Acc=0.9475]
Evaluating: 100%|      | 100/100 [00:02<00:00, 48.13it/s, Loss=0.1356,
IoU=0.6406, Acc=0.9475]

Train - Loss: 0.0969, IoU: 0.7995, Acc: 0.9660
Val   - Loss: 0.4446, IoU: 0.5843, Acc: 0.8779
Best Val IoU: 0.5951
```

Epoch 19/30

```
Training: 100%|      | 400/400 [00:14<00:00, 27.64it/s, Loss=0.2063,
IoU=0.7578, Acc=0.9240]
Training: 100%|      | 400/400 [00:14<00:00, 27.64it/s, Loss=0.2063,
IoU=0.7578, Acc=0.9240]
Evaluating: 100%|      | 100/100 [00:02<00:00, 45.07it/s, Loss=0.1255,
IoU=0.6523, Acc=0.9529]
Evaluating: 100%|      | 100/100 [00:02<00:00, 45.07it/s, Loss=0.1255,
IoU=0.6523, Acc=0.9529]

Train - Loss: 0.0980, IoU: 0.8034, Acc: 0.9654
Val   - Loss: 0.4421, IoU: 0.5876, Acc: 0.8771
Best Val IoU: 0.5951
```

Epoch 20/30

```
Training: 100%|      | 400/400 [00:15<00:00, 25.87it/s, Loss=0.0271,
IoU=0.9022, Acc=0.9908]
Training: 100%|      | 400/400 [00:15<00:00, 25.87it/s, Loss=0.0271,
IoU=0.9022, Acc=0.9908]
Evaluating: 100%|      | 100/100 [00:02<00:00, 44.29it/s, Loss=0.1496,
IoU=0.6296, Acc=0.9420]
Evaluating: 100%|      | 100/100 [00:02<00:00, 44.29it/s, Loss=0.1496,
```

IoU=0.6296, Acc=0.9420]

Train - Loss: 0.0922, IoU: 0.8014, Acc: 0.9674

Val - Loss: 0.4528, IoU: 0.5776, Acc: 0.8763

Best Val IoU: 0.5951

Epoch 21/30

Training: 100% | 400/400 [00:14<00:00, 27.51it/s, Loss=0.0573, IoU=0.8562, Acc=0.9762]

Training: 100% | 400/400 [00:14<00:00, 27.51it/s, Loss=0.0573, IoU=0.8562, Acc=0.9762]

Evaluating: 100% | 100/100 [00:02<00:00, 48.03it/s, Loss=0.1384, IoU=0.6388, Acc=0.9466]

Evaluating: 100% | 100/100 [00:02<00:00, 48.03it/s, Loss=0.1384, IoU=0.6388, Acc=0.9466]

Train - Loss: 0.0924, IoU: 0.8034, Acc: 0.9672

Val - Loss: 0.4387, IoU: 0.5811, Acc: 0.8778

Best Val IoU: 0.5951

Epoch 22/30

Training: 100% | 400/400 [00:14<00:00, 27.57it/s, Loss=0.0678, IoU=0.6473, Acc=0.9770]

Training: 100% | 400/400 [00:14<00:00, 27.57it/s, Loss=0.0678, IoU=0.6473, Acc=0.9770]

Evaluating: 100% | 100/100 [00:02<00:00, 45.95it/s, Loss=0.1303, IoU=0.6458, Acc=0.9501]

Evaluating: 100% | 100/100 [00:02<00:00, 45.95it/s, Loss=0.1303, IoU=0.6458, Acc=0.9501]

Train - Loss: 0.0915, IoU: 0.8068, Acc: 0.9673

Val - Loss: 0.4437, IoU: 0.5804, Acc: 0.8787

Best Val IoU: 0.5951

Epoch 23/30

Training: 100% | 400/400 [00:15<00:00, 26.32it/s, Loss=0.0558, IoU=0.9124, Acc=0.9816]

Training: 100% | 400/400 [00:15<00:00, 26.32it/s, Loss=0.0558, IoU=0.9124, Acc=0.9816]

Evaluating: 100% | 100/100 [00:02<00:00, 47.47it/s, Loss=0.1283, IoU=0.6507, Acc=0.9522]

Evaluating: 100% | 100/100 [00:02<00:00, 47.47it/s, Loss=0.1283, IoU=0.6507, Acc=0.9522]

Train - Loss: 0.0899, IoU: 0.8103, Acc: 0.9679

Val - Loss: 0.4417, IoU: 0.5840, Acc: 0.8783

Best Val IoU: 0.5951

Epoch 24/30

```
Training: 100%|     | 400/400 [00:14<00:00, 27.61it/s, Loss=0.0463,
IoU=0.9263, Acc=0.9834]
Training: 100%|     | 400/400 [00:14<00:00, 27.61it/s, Loss=0.0463,
IoU=0.9263, Acc=0.9834]
Evaluating: 100%|    | 100/100 [00:02<00:00, 49.04it/s, Loss=0.1287,
IoU=0.6466, Acc=0.9507]
Evaluating: 100%|    | 100/100 [00:02<00:00, 49.04it/s, Loss=0.1287,
IoU=0.6466, Acc=0.9507]

Train - Loss: 0.0882, IoU: 0.8096, Acc: 0.9684
Val - Loss: 0.4333, IoU: 0.5836, Acc: 0.8783
Best Val IoU: 0.5951
```

Epoch 25/30

```
Training: 100%|     | 400/400 [00:14<00:00, 27.50it/s, Loss=0.1833,
IoU=0.7809, Acc=0.9309]
Training: 100%|     | 400/400 [00:14<00:00, 27.50it/s, Loss=0.1833,
IoU=0.7809, Acc=0.9309]
Evaluating: 100%|    | 100/100 [00:02<00:00, 47.98it/s, Loss=0.1200,
IoU=0.6563, Acc=0.9548]
Evaluating: 100%|    | 100/100 [00:02<00:00, 47.98it/s, Loss=0.1200,
IoU=0.6563, Acc=0.9548]

Train - Loss: 0.0884, IoU: 0.8065, Acc: 0.9681
Val - Loss: 0.4423, IoU: 0.5841, Acc: 0.8788
Best Val IoU: 0.5951
```

Epoch 26/30

```
Training: 100%|     | 400/400 [00:14<00:00, 27.55it/s, Loss=0.1654,
IoU=0.6634, Acc=0.9399]
Training: 100%|     | 400/400 [00:14<00:00, 27.55it/s, Loss=0.1654,
IoU=0.6634, Acc=0.9399]
Evaluating: 100%|    | 100/100 [00:02<00:00, 48.30it/s, Loss=0.1236,
IoU=0.6523, Acc=0.9534]
Evaluating: 100%|    | 100/100 [00:02<00:00, 48.30it/s, Loss=0.1236,
IoU=0.6523, Acc=0.9534]

Train - Loss: 0.0872, IoU: 0.8079, Acc: 0.9686
Val - Loss: 0.4392, IoU: 0.5781, Acc: 0.8780
Best Val IoU: 0.5951
```

Epoch 27/30

```
Training: 100%|     | 400/400 [00:14<00:00, 27.51it/s, Loss=0.0996,
IoU=0.8868, Acc=0.9646]
Training: 100%|     | 400/400 [00:14<00:00, 27.51it/s, Loss=0.0996,
IoU=0.8868, Acc=0.9646]
```

```
Evaluating: 100%|      | 100/100 [00:02<00:00, 48.54it/s, Loss=0.1363,  
IoU=0.6404, Acc=0.9472]  
Evaluating: 100%|      | 100/100 [00:02<00:00, 48.54it/s, Loss=0.1363,  
IoU=0.6404, Acc=0.9472]  
Train - Loss: 0.0869, IoU: 0.8097, Acc: 0.9687  
Val - Loss: 0.4508, IoU: 0.5874, Acc: 0.8800  
Best Val IoU: 0.5951
```

Epoch 28/30

```
Training: 100%|      | 400/400 [00:14<00:00, 27.34it/s, Loss=0.0372,  
IoU=0.9535, Acc=0.9860]  
Training: 100%|      | 400/400 [00:14<00:00, 27.34it/s, Loss=0.0372,  
IoU=0.9535, Acc=0.9860]  
Evaluating: 100%|      | 100/100 [00:02<00:00, 43.74it/s, Loss=0.1656,  
IoU=0.6251, Acc=0.9384]  
Evaluating: 100%|      | 100/100 [00:02<00:00, 43.74it/s, Loss=0.1656,  
IoU=0.6251, Acc=0.9384]  
Train - Loss: 0.0894, IoU: 0.8112, Acc: 0.9679  
Val - Loss: 0.4452, IoU: 0.5863, Acc: 0.8791  
Best Val IoU: 0.5951
```

Epoch 29/30

```
Training: 100%|      | 400/400 [00:15<00:00, 26.02it/s, Loss=0.1001,  
IoU=0.7039, Acc=0.9605]  
Training: 100%|      | 400/400 [00:15<00:00, 26.02it/s, Loss=0.1001,  
IoU=0.7039, Acc=0.9605]  
Evaluating: 100%|      | 100/100 [00:02<00:00, 46.43it/s, Loss=0.1446,  
IoU=0.6286, Acc=0.9414]  
Evaluating: 100%|      | 100/100 [00:02<00:00, 46.43it/s, Loss=0.1446,  
IoU=0.6286, Acc=0.9414]  
Train - Loss: 0.0896, IoU: 0.8059, Acc: 0.9680  
Val - Loss: 0.4525, IoU: 0.5761, Acc: 0.8768  
Best Val IoU: 0.5951
```

Epoch 30/30

```
Training: 100%|      | 400/400 [00:14<00:00, 27.31it/s, Loss=0.1145,  
IoU=0.8924, Acc=0.9568]  
Training: 100%|      | 400/400 [00:14<00:00, 27.31it/s, Loss=0.1145,  
IoU=0.8924, Acc=0.9568]  
Evaluating: 100%|      | 100/100 [00:02<00:00, 48.79it/s, Loss=0.1351,  
IoU=0.6449, Acc=0.9496]
```

```
Train - Loss: 0.0865, IoU: 0.8093, Acc: 0.9689  
Val - Loss: 0.4533, IoU: 0.5819, Acc: 0.8782
```

Best Val IoU: 0.5951

Training completed! Best validation IoU: 0.5951

Transpose convolution training completed!

1.5.2 2. FCN-8s with Bilinear Interpolation

```
[12]: # Create and train FCN with bilinear interpolation
print("==" * 80)
print("TRAINING FCN-8s WITH BILINEAR INTERPOLATION")
print("==" * 80)

# Initialize model
fcn_bilinear = FCN8s(num_classes=21, upsampling='bilinear').to(device)

# Check if pretrained model exists
bilinear_model_path = 'models/task_2/fcn8s_bilinear_best.pth'
if os.path.exists(bilinear_model_path):
    print("Loading pretrained bilinear interpolation model...")
    fcn_bilinear.load_state_dict(torch.load(bilinear_model_path, map_location=device))

# Quick evaluation
val_loss, val_iou, val_acc = evaluate_model(fcn_bilinear, test_loader, nn.CrossEntropyLoss(ignore_index=255), device)
print(f"Loaded model performance - Loss: {val_loss:.4f}, IoU: {val_iou:.4f}, Acc: {val_acc:.4f}")

# Create dummy history for consistency
bilinear_history = {
    'train_loss': [], 'train_iou': [], 'train_accuracy': [],
    'val_loss': [], 'val_iou': [], 'val_accuracy': []
}
else:
    print("Training new bilinear interpolation model...")
    bilinear_history = train_model(fcn_bilinear, train_loader, test_loader, num_epochs=30, learning_rate=1e-4)

print("\nBilinear interpolation training completed!")
```

```
=====
TRAINING FCN-8s WITH BILINEAR INTERPOLATION
=====
Training new bilinear interpolation model...
Starting training for 30 epochs...
Model: FCN8s with bilinear upsampling
```

```
-----  
Epoch 1/30  
Training new bilinear interpolation model...  
Starting training for 30 epochs..  
Model: FCN8s with bilinear upsampling  
-----  
  
Epoch 1/30  
  
Training: 100%|      | 400/400 [00:14<00:00, 27.78it/s, Loss=1.2315,  
IoU=0.3315, Acc=0.7023]  
Training: 100%|      | 400/400 [00:14<00:00, 27.78it/s, Loss=1.2315,  
IoU=0.3315, Acc=0.7023]  
Evaluating: 100%|      | 100/100 [00:02<00:00, 49.73it/s, Loss=0.3661,  
IoU=0.5535, Acc=0.9008]  
Evaluating: 100%|      | 100/100 [00:02<00:00, 49.73it/s, Loss=0.3661,  
IoU=0.5535, Acc=0.9008]  
  
Train - Loss: 1.0396, IoU: 0.3154, Acc: 0.7571  
Val - Loss: 0.6264, IoU: 0.4667, Acc: 0.8344  
Best Val IoU: 0.4667  
  
Epoch 2/30  
  
Training: 100%|      | 400/400 [00:14<00:00, 27.99it/s, Loss=1.0719,  
IoU=0.5435, Acc=0.7550]  
Training: 100%|      | 400/400 [00:14<00:00, 27.99it/s, Loss=1.0719,  
IoU=0.5435, Acc=0.7550]  
Evaluating: 100%|      | 100/100 [00:02<00:00, 48.84it/s, Loss=0.2962,  
IoU=0.5861, Acc=0.9192]  
  
Train - Loss: 0.6287, IoU: 0.4813, Acc: 0.8337  
Val - Loss: 0.5179, IoU: 0.5320, Acc: 0.8577  
Best Val IoU: 0.5320  
  
Epoch 3/30  
  
Training: 100%|      | 400/400 [00:14<00:00, 27.80it/s, Loss=0.3585,  
IoU=0.5235, Acc=0.9445]  
Training: 100%|      | 400/400 [00:14<00:00, 27.80it/s, Loss=0.3585,  
IoU=0.5235, Acc=0.9445]  
Evaluating: 100%|      | 100/100 [00:02<00:00, 49.91it/s, Loss=0.1778,  
IoU=0.6735, Acc=0.9594]  
  
Train - Loss: 0.4765, IoU: 0.5638, Acc: 0.8697  
Val - Loss: 0.4823, IoU: 0.5735, Acc: 0.8674  
Best Val IoU: 0.5735
```

Epoch 4/30

```
Training: 100%|     | 400/400 [00:14<00:00, 27.87it/s, Loss=1.8873,
IoU=0.3481, Acc=0.3748]
Training: 100%|     | 400/400 [00:14<00:00, 27.87it/s, Loss=1.8873,
IoU=0.3481, Acc=0.3748]
Evaluating: 100%|    | 100/100 [00:02<00:00, 48.27it/s, Loss=0.1605,
IoU=0.6629, Acc=0.9558]
Evaluating: 100%|    | 100/100 [00:02<00:00, 48.27it/s, Loss=0.1605,
IoU=0.6629, Acc=0.9558]

Train - Loss: 0.3754, IoU: 0.6203, Acc: 0.8938
Val   - Loss: 0.4258, IoU: 0.5875, Acc: 0.8794
Best Val IoU: 0.5875
```

Epoch 5/30

```
Training: 100%|     | 400/400 [00:14<00:00, 27.81it/s, Loss=0.1127,
IoU=0.8640, Acc=0.9767]
Training: 100%|     | 400/400 [00:14<00:00, 27.81it/s, Loss=0.1127,
IoU=0.8640, Acc=0.9767]
Evaluating: 100%|    | 100/100 [00:02<00:00, 49.88it/s, Loss=0.1247,
IoU=0.6808, Acc=0.9641]

Train - Loss: 0.2886, IoU: 0.6701, Acc: 0.9166
Val   - Loss: 0.4033, IoU: 0.6008, Acc: 0.8829
Best Val IoU: 0.6008
```

Epoch 6/30

```
Training: 100%|     | 400/400 [00:14<00:00, 27.84it/s, Loss=0.0507,
IoU=0.7972, Acc=0.9937]
Training: 100%|     | 400/400 [00:14<00:00, 27.84it/s, Loss=0.0507,
IoU=0.7972, Acc=0.9937]
Evaluating: 100%|    | 100/100 [00:02<00:00, 49.67it/s, Loss=0.1357,
IoU=0.6770, Acc=0.9634]
Evaluating: 100%|    | 100/100 [00:02<00:00, 49.67it/s, Loss=0.1357,
IoU=0.6770, Acc=0.9634]

Train - Loss: 0.2352, IoU: 0.7007, Acc: 0.9309
Val   - Loss: 0.4167, IoU: 0.5952, Acc: 0.8790
Best Val IoU: 0.6008
```

Epoch 7/30

```
Training: 100%|     | 400/400 [00:14<00:00, 28.03it/s, Loss=0.1095,
IoU=0.5497, Acc=0.9638]
Training: 100%|     | 400/400 [00:14<00:00, 28.03it/s, Loss=0.1095,
IoU=0.5497, Acc=0.9638]
Evaluating: 100%|    | 100/100 [00:01<00:00, 50.15it/s, Loss=0.1483,
```

```
[IoU=0.6418, Acc=0.9463]  
Evaluating: 100% | 100/100 [00:01<00:00, 50.15it/s, Loss=0.1483,  
IoU=0.6418, Acc=0.9463]
```

```
Train - Loss: 0.2039, IoU: 0.7228, Acc: 0.9379  
Val - Loss: 0.4235, IoU: 0.5983, Acc: 0.8769  
Best Val IoU: 0.6008
```

Epoch 8/30

```
Training: 100% | 400/400 [00:14<00:00, 28.08it/s, Loss=0.3194,  
IoU=0.7957, Acc=0.8993]  
Training: 100% | 400/400 [00:14<00:00, 28.08it/s, Loss=0.3194,  
IoU=0.7957, Acc=0.8993]  
Evaluating: 100% | 100/100 [00:01<00:00, 50.94it/s, Loss=0.1170,  
IoU=0.6783, Acc=0.9631]  
Evaluating: 100% | 100/100 [00:01<00:00, 50.94it/s, Loss=0.1170,  
IoU=0.6783, Acc=0.9631]
```

```
Train - Loss: 0.1867, IoU: 0.7334, Acc: 0.9429  
Val - Loss: 0.4221, IoU: 0.6008, Acc: 0.8807  
Best Val IoU: 0.6008
```

Epoch 9/30

```
Training: 100% | 400/400 [00:14<00:00, 28.11it/s, Loss=0.0645,  
IoU=0.8736, Acc=0.9874]  
Training: 100% | 400/400 [00:14<00:00, 28.11it/s, Loss=0.0645,  
IoU=0.8736, Acc=0.9874]  
Evaluating: 100% | 100/100 [00:02<00:00, 49.33it/s, Loss=0.1210,  
IoU=0.6686, Acc=0.9575]
```

```
Train - Loss: 0.1507, IoU: 0.7593, Acc: 0.9517  
Val - Loss: 0.4185, IoU: 0.6127, Acc: 0.8786  
Best Val IoU: 0.6127
```

Epoch 10/30

```
Training: 100% | 400/400 [00:14<00:00, 27.95it/s, Loss=0.0605,  
IoU=0.7010, Acc=0.9816]  
Training: 100% | 400/400 [00:14<00:00, 27.95it/s, Loss=0.0605,  
IoU=0.7010, Acc=0.9816]  
Evaluating: 100% | 100/100 [00:02<00:00, 49.94it/s, Loss=0.1283,  
IoU=0.6761, Acc=0.9607]  
Evaluating: 100% | 100/100 [00:02<00:00, 49.94it/s, Loss=0.1283,  
IoU=0.6761, Acc=0.9607]
```

```
Train - Loss: 0.1408, IoU: 0.7627, Acc: 0.9543  
Val - Loss: 0.4257, IoU: 0.6014, Acc: 0.8769  
Best Val IoU: 0.6127
```

Epoch 11/30

```
Training: 100%|      | 400/400 [00:14<00:00, 28.06it/s, Loss=0.0420,
IoU=0.7360, Acc=0.9914]
Training: 100%|      | 400/400 [00:14<00:00, 28.06it/s, Loss=0.0420,
IoU=0.7360, Acc=0.9914]
Evaluating: 100%|     | 100/100 [00:01<00:00, 50.26it/s, Loss=0.1246,
IoU=0.6748, Acc=0.9604]
Evaluating: 100%|     | 100/100 [00:01<00:00, 50.26it/s, Loss=0.1246,
IoU=0.6748, Acc=0.9604]

Train - Loss: 0.1216, IoU: 0.7811, Acc: 0.9597
Val   - Loss: 0.4216, IoU: 0.5970, Acc: 0.8808
Best Val IoU: 0.6127
```

Epoch 12/30

```
Training: 100%|      | 400/400 [00:14<00:00, 27.87it/s, Loss=0.0986,
IoU=0.6844, Acc=0.9624]
Training: 100%|      | 400/400 [00:14<00:00, 27.87it/s, Loss=0.0986,
IoU=0.6844, Acc=0.9624]
Evaluating: 100%|     | 100/100 [00:02<00:00, 47.95it/s, Loss=0.1105,
IoU=0.6789, Acc=0.9629]
Evaluating: 100%|     | 100/100 [00:02<00:00, 47.95it/s, Loss=0.1105,
IoU=0.6789, Acc=0.9629]

Train - Loss: 0.1152, IoU: 0.7837, Acc: 0.9608
Val   - Loss: 0.4059, IoU: 0.6101, Acc: 0.8851
Best Val IoU: 0.6127
```

Epoch 13/30

```
Training: 100%|      | 400/400 [00:14<00:00, 27.86it/s, Loss=0.0694,
IoU=0.8278, Acc=0.9759]
Training: 100%|      | 400/400 [00:14<00:00, 27.86it/s, Loss=0.0694,
IoU=0.8278, Acc=0.9759]
Evaluating: 100%|     | 100/100 [00:01<00:00, 50.26it/s, Loss=0.1108,
IoU=0.6802, Acc=0.9638]
Evaluating: 100%|     | 100/100 [00:01<00:00, 50.26it/s, Loss=0.1108,
IoU=0.6802, Acc=0.9638]

Train - Loss: 0.1156, IoU: 0.7854, Acc: 0.9609
Val   - Loss: 0.4013, IoU: 0.6107, Acc: 0.8845
Best Val IoU: 0.6127
```

Epoch 14/30

```
Training: 100%|      | 400/400 [00:14<00:00, 27.79it/s, Loss=0.0785,
IoU=0.7656, Acc=0.9727]
Training: 100%|      | 400/400 [00:14<00:00, 27.79it/s, Loss=0.0785,
```

```
[IoU=0.7656, Acc=0.9727]
Evaluating: 100%|      | 100/100 [00:02<00:00, 48.51it/s, Loss=0.1081,
IoU=0.6784, Acc=0.9630]
Evaluating: 100%|      | 100/100 [00:02<00:00, 48.51it/s, Loss=0.1081,
IoU=0.6784, Acc=0.9630]

Train - Loss: 0.1060, IoU: 0.7929, Acc: 0.9637
Val   - Loss: 0.4032, IoU: 0.6093, Acc: 0.8863
Best Val IoU: 0.6127
```

Epoch 15/30

```
Training: 100%|      | 400/400 [00:14<00:00, 27.96it/s, Loss=0.0944,
IoU=0.9287, Acc=0.9704]
Training: 100%|      | 400/400 [00:14<00:00, 27.96it/s, Loss=0.0944,
IoU=0.9287, Acc=0.9704]
Evaluating: 100%|      | 100/100 [00:02<00:00, 49.95it/s, Loss=0.1254,
IoU=0.6724, Acc=0.9591]
Evaluating: 100%|      | 100/100 [00:02<00:00, 49.95it/s, Loss=0.1254,
IoU=0.6724, Acc=0.9591]

Train - Loss: 0.1041, IoU: 0.7951, Acc: 0.9643
Val   - Loss: 0.4160, IoU: 0.6065, Acc: 0.8814
Best Val IoU: 0.6127
```

Epoch 16/30

```
Training: 100%|      | 400/400 [00:14<00:00, 27.86it/s, Loss=0.0543,
IoU=0.9456, Acc=0.9831]
Training: 100%|      | 400/400 [00:14<00:00, 27.86it/s, Loss=0.0543,
IoU=0.9456, Acc=0.9831]
Evaluating: 100%|      | 100/100 [00:02<00:00, 49.32it/s, Loss=0.1175,
IoU=0.6679, Acc=0.9589]
Evaluating: 100%|      | 100/100 [00:02<00:00, 49.32it/s, Loss=0.1175,
IoU=0.6679, Acc=0.9589]

Train - Loss: 0.1006, IoU: 0.7968, Acc: 0.9650
Val   - Loss: 0.4152, IoU: 0.6088, Acc: 0.8847
Best Val IoU: 0.6127
```

Epoch 17/30

```
Training: 100%|      | 400/400 [00:14<00:00, 27.89it/s, Loss=0.1446,
IoU=0.8692, Acc=0.9480]
Training: 100%|      | 400/400 [00:14<00:00, 27.89it/s, Loss=0.1446,
IoU=0.8692, Acc=0.9480]
Evaluating: 100%|      | 100/100 [00:02<00:00, 47.63it/s, Loss=0.1107,
IoU=0.6766, Acc=0.9618]
Evaluating: 100%|      | 100/100 [00:02<00:00, 47.63it/s, Loss=0.1107,
IoU=0.6766, Acc=0.9618]
```

```
Train - Loss: 0.0961, IoU: 0.8033, Acc: 0.9663
Val   - Loss: 0.4181, IoU: 0.6037, Acc: 0.8837
Best Val IoU: 0.6127
```

Epoch 18/30

```
Training: 100%|      | 400/400 [00:14<00:00, 27.96it/s, Loss=0.0601,
IoU=0.9188, Acc=0.9778]
Training: 100%|      | 400/400 [00:14<00:00, 27.96it/s, Loss=0.0601,
IoU=0.9188, Acc=0.9778]
Evaluating: 100%|     | 100/100 [00:02<00:00, 47.20it/s, Loss=0.1147,
IoU=0.6759, Acc=0.9608]
```

```
Train - Loss: 0.0989, IoU: 0.8032, Acc: 0.9654
Val   - Loss: 0.4082, IoU: 0.6136, Acc: 0.8864
Best Val IoU: 0.6136
```

Epoch 19/30

```
Training: 100%|      | 400/400 [00:14<00:00, 27.74it/s, Loss=0.0833,
IoU=0.5977, Acc=0.9677]
Training: 100%|      | 400/400 [00:14<00:00, 27.74it/s, Loss=0.0833,
IoU=0.5977, Acc=0.9677]
Evaluating: 100%|     | 100/100 [00:02<00:00, 49.42it/s, Loss=0.1088,
IoU=0.6792, Acc=0.9633]
Evaluating: 100%|     | 100/100 [00:02<00:00, 49.42it/s, Loss=0.1088,
IoU=0.6792, Acc=0.9633]
```

```
Train - Loss: 0.0928, IoU: 0.8076, Acc: 0.9674
Val   - Loss: 0.4166, IoU: 0.6060, Acc: 0.8848
Best Val IoU: 0.6136
```

Epoch 20/30

```
Training: 100%|      | 400/400 [00:14<00:00, 27.92it/s, Loss=0.0419,
IoU=0.8616, Acc=0.9861]
Training: 100%|      | 400/400 [00:14<00:00, 27.92it/s, Loss=0.0419,
IoU=0.8616, Acc=0.9861]
Evaluating: 100%|     | 100/100 [00:02<00:00, 49.66it/s, Loss=0.1198,
IoU=0.6659, Acc=0.9577]
Evaluating: 100%|     | 100/100 [00:02<00:00, 49.66it/s, Loss=0.1198,
IoU=0.6659, Acc=0.9577]
```

```
Train - Loss: 0.0926, IoU: 0.8050, Acc: 0.9669
Val   - Loss: 0.4200, IoU: 0.6092, Acc: 0.8842
Best Val IoU: 0.6136
```

Epoch 21/30

```
Training: 100%|      | 400/400 [00:14<00:00, 27.95it/s, Loss=0.1420,
```

```
[IoU=0.8335, Acc=0.9554]
Training: 100%|      | 400/400 [00:14<00:00, 27.95it/s, Loss=0.1420,
IoU=0.8335, Acc=0.9554]
Evaluating: 100%|      | 100/100 [00:02<00:00, 48.67it/s, Loss=0.1195,
IoU=0.6748, Acc=0.9606]
Evaluating: 100%|      | 100/100 [00:02<00:00, 48.67it/s, Loss=0.1195,
IoU=0.6748, Acc=0.9606]

Train - Loss: 0.0908, IoU: 0.8118, Acc: 0.9681
Val   - Loss: 0.4257, IoU: 0.6041, Acc: 0.8822
Best Val IoU: 0.6136
```

Epoch 22/30

```
Training: 100%|      | 400/400 [00:14<00:00, 28.12it/s, Loss=0.0589,
IoU=0.7799, Acc=0.9751]
Training: 100%|      | 400/400 [00:14<00:00, 28.12it/s, Loss=0.0589,
IoU=0.7799, Acc=0.9751]
Evaluating: 100%|      | 100/100 [00:02<00:00, 48.30it/s, Loss=0.1155,
IoU=0.6741, Acc=0.9607]
Evaluating: 100%|      | 100/100 [00:02<00:00, 48.30it/s, Loss=0.1155,
IoU=0.6741, Acc=0.9607]

Train - Loss: 0.0904, IoU: 0.8111, Acc: 0.9680
Val   - Loss: 0.4201, IoU: 0.6068, Acc: 0.8851
Best Val IoU: 0.6136
```

Epoch 23/30

```
Training: 100%|      | 400/400 [00:14<00:00, 27.81it/s, Loss=0.3386,
IoU=0.7739, Acc=0.8733]
Training: 100%|      | 400/400 [00:14<00:00, 27.81it/s, Loss=0.3386,
IoU=0.7739, Acc=0.8733]
Evaluating: 100%|      | 100/100 [00:02<00:00, 46.04it/s, Loss=0.1156,
IoU=0.6700, Acc=0.9603]
Evaluating: 100%|      | 100/100 [00:02<00:00, 46.04it/s, Loss=0.1156,
IoU=0.6700, Acc=0.9603]

Train - Loss: 0.0886, IoU: 0.8084, Acc: 0.9684
Val   - Loss: 0.4277, IoU: 0.6027, Acc: 0.8838
Best Val IoU: 0.6136
```

Epoch 24/30

```
Training: 100%|      | 400/400 [00:14<00:00, 27.60it/s, Loss=0.0638,
IoU=0.6680, Acc=0.9739]
Training: 100%|      | 400/400 [00:14<00:00, 27.60it/s, Loss=0.0638,
IoU=0.6680, Acc=0.9739]
Evaluating: 100%|      | 100/100 [00:02<00:00, 48.99it/s, Loss=0.1116,
IoU=0.6759, Acc=0.9616]
Evaluating: 100%|      | 100/100 [00:02<00:00, 48.99it/s, Loss=0.1116,
```

IoU=0.6759, Acc=0.9616]

Train - Loss: 0.0889, IoU: 0.8091, Acc: 0.9683

Val - Loss: 0.4148, IoU: 0.6081, Acc: 0.8874

Best Val IoU: 0.6136

Epoch 25/30

Training: 100% | 400/400 [00:14<00:00, 27.91it/s, Loss=0.0875, IoU=0.8824, Acc=0.9713]

Training: 100% | 400/400 [00:14<00:00, 27.91it/s, Loss=0.0875, IoU=0.8824, Acc=0.9713]

Evaluating: 100% | 100/100 [00:02<00:00, 48.23it/s, Loss=0.1150, IoU=0.6698, Acc=0.9592]

Evaluating: 100% | 100/100 [00:02<00:00, 48.23it/s, Loss=0.1150, IoU=0.6698, Acc=0.9592]

Train - Loss: 0.0892, IoU: 0.8101, Acc: 0.9681

Val - Loss: 0.4243, IoU: 0.6052, Acc: 0.8854

Best Val IoU: 0.6136

Epoch 26/30

Training: 100% | 400/400 [00:14<00:00, 27.86it/s, Loss=0.0949, IoU=0.6800, Acc=0.9658]

Training: 100% | 400/400 [00:14<00:00, 27.86it/s, Loss=0.0949, IoU=0.6800, Acc=0.9658]

Evaluating: 100% | 100/100 [00:02<00:00, 49.76it/s, Loss=0.1063, IoU=0.6762, Acc=0.9633]

Evaluating: 100% | 100/100 [00:02<00:00, 49.76it/s, Loss=0.1063, IoU=0.6762, Acc=0.9633]

Train - Loss: 0.0877, IoU: 0.8125, Acc: 0.9687

Val - Loss: 0.4245, IoU: 0.6014, Acc: 0.8842

Best Val IoU: 0.6136

Epoch 27/30

Training: 100% | 400/400 [00:14<00:00, 28.17it/s, Loss=0.0475, IoU=0.9644, Acc=0.9840]

Training: 100% | 400/400 [00:14<00:00, 28.17it/s, Loss=0.0475, IoU=0.9644, Acc=0.9840]

Evaluating: 100% | 100/100 [00:02<00:00, 49.91it/s, Loss=0.1127, IoU=0.6725, Acc=0.9602]

Evaluating: 100% | 100/100 [00:02<00:00, 49.91it/s, Loss=0.1127, IoU=0.6725, Acc=0.9602]

Train - Loss: 0.0900, IoU: 0.8073, Acc: 0.9678

Val - Loss: 0.4135, IoU: 0.6100, Acc: 0.8858

Best Val IoU: 0.6136

Epoch 28/30

```
Training: 100%|     | 400/400 [00:14<00:00, 28.18it/s, Loss=0.0611,
IoU=0.8877, Acc=0.9751]
Training: 100%|     | 400/400 [00:14<00:00, 28.18it/s, Loss=0.0611,
IoU=0.8877, Acc=0.9751]
Evaluating: 100%|     | 100/100 [00:01<00:00, 50.07it/s, Loss=0.1138,
IoU=0.6729, Acc=0.9600]
Evaluating: 100%|     | 100/100 [00:01<00:00, 50.07it/s, Loss=0.1138,
IoU=0.6729, Acc=0.9600]

Train - Loss: 0.0886, IoU: 0.8119, Acc: 0.9686
Val - Loss: 0.4173, IoU: 0.6070, Acc: 0.8847
Best Val IoU: 0.6136
```

Epoch 29/30

```
Training: 100%|     | 400/400 [00:14<00:00, 28.14it/s, Loss=0.1300,
IoU=0.8675, Acc=0.9551]
Training: 100%|     | 400/400 [00:14<00:00, 28.14it/s, Loss=0.1300,
IoU=0.8675, Acc=0.9551]
Evaluating: 100%|     | 100/100 [00:02<00:00, 49.74it/s, Loss=0.1158,
IoU=0.6706, Acc=0.9591]
Evaluating: 100%|     | 100/100 [00:02<00:00, 49.74it/s, Loss=0.1158,
IoU=0.6706, Acc=0.9591]

Train - Loss: 0.0881, IoU: 0.8149, Acc: 0.9686
Val - Loss: 0.4222, IoU: 0.6063, Acc: 0.8839
Best Val IoU: 0.6136
```

Epoch 30/30

```
Training: 100%|     | 400/400 [00:14<00:00, 28.24it/s, Loss=0.0850,
IoU=0.9127, Acc=0.9720]
Training: 100%|     | 400/400 [00:14<00:00, 28.24it/s, Loss=0.0850,
IoU=0.9127, Acc=0.9720]
Evaluating: 100%|     | 100/100 [00:01<00:00, 50.42it/s, Loss=0.1064,
IoU=0.6782, Acc=0.9632]

Train - Loss: 0.0867, IoU: 0.8110, Acc: 0.9689
Val - Loss: 0.4271, IoU: 0.6040, Acc: 0.8835
Best Val IoU: 0.6136
```

Training completed! Best validation IoU: 0.6136

Bilinear interpolation training completed!

1.6 Visualization and Analysis

1.6.1 1. Training Curves

```
[13]: # Plot training curves for both models
def plot_training_curves(transpose_history, bilinear_history):
    """Plot training curves comparing both models"""
    fig, axes = plt.subplots(2, 3, figsize=(18, 12))

    # If histories are empty (loaded pretrained models), skip plotting
    if not transpose_history['train_loss'] and not bilinear_history['train_loss']:
        print("No training history available (models were loaded from checkpoints)")
        return

    epochs = range(1, len(transpose_history['train_loss']) + 1) if transpose_history['train_loss'] else []
    epochs_b = range(1, len(bilinear_history['train_loss']) + 1) if bilinear_history['train_loss'] else []

    # Loss curves
    if epochs:
        axes[0, 0].plot(epochs, transpose_history['train_loss'], 'b-', label='Transpose - Train', linewidth=2)
        axes[0, 0].plot(epochs, transpose_history['val_loss'], 'b--', label='Transpose - Val', linewidth=2)
        if epochs_b:
            axes[0, 0].plot(epochs_b, bilinear_history['train_loss'], 'r-', label='Bilinear - Train', linewidth=2)
            axes[0, 0].plot(epochs_b, bilinear_history['val_loss'], 'r--', label='Bilinear - Val', linewidth=2)
        axes[0, 0].set_title('Training & Validation Loss', fontsize=14, fontweight='bold')
        axes[0, 0].set_xlabel('Epoch')
        axes[0, 0].set_ylabel('Loss')
        axes[0, 0].legend()
        axes[0, 0].grid(True, alpha=0.3)

    # IoU curves
    if epochs:
        axes[0, 1].plot(epochs, transpose_history['train_iou'], 'b-', label='Transpose - Train', linewidth=2)
        axes[0, 1].plot(epochs, transpose_history['val_iou'], 'b--', label='Transpose - Val', linewidth=2)
    if epochs_b:
        axes[0, 1].plot(epochs_b, bilinear_history['train_iou'], 'r-', label='Bilinear - Train', linewidth=2)
        axes[0, 1].plot(epochs_b, bilinear_history['val_iou'], 'r--', label='Bilinear - Val', linewidth=2)
```

```

        axes[0, 1].plot(epochs_b, bilinear_history['train_iou'], 'r-',  

    ↪label='Bilinear - Train', linewidth=2)  

        axes[0, 1].plot(epochs_b, bilinear_history['val_iou'], 'r--',  

    ↪label='Bilinear - Val', linewidth=2)  

        axes[0, 1].set_title('Training & Validation IoU', fontsize=14,  

    ↪fontweight='bold')  

        axes[0, 1].set_xlabel('Epoch')  

        axes[0, 1].set_ylabel('Mean IoU')  

        axes[0, 1].legend()  

        axes[0, 1].grid(True, alpha=0.3)

# Accuracy curves
if epochs:  

    axes[0, 2].plot(epochs, transpose_history['train_accuracy'], 'b-',  

    ↪label='Transpose - Train', linewidth=2)  

    axes[0, 2].plot(epochs, transpose_history['val_accuracy'], 'b--',  

    ↪label='Transpose - Val', linewidth=2)  

    if epochs_b:  

        axes[0, 2].plot(epochs_b, bilinear_history['train_accuracy'], 'r-',  

    ↪label='Bilinear - Train', linewidth=2)  

        axes[0, 2].plot(epochs_b, bilinear_history['val_accuracy'], 'r--',  

    ↪label='Bilinear - Val', linewidth=2)  

        axes[0, 2].set_title('Training & Validation Accuracy', fontsize=14,  

    ↪fontweight='bold')  

        axes[0, 2].set_xlabel('Epoch')  

        axes[0, 2].set_ylabel('Pixel Accuracy')  

        axes[0, 2].legend()  

        axes[0, 2].grid(True, alpha=0.3)

# Individual loss comparison
if epochs:  

    axes[1, 0].plot(epochs, transpose_history['val_loss'], 'b-',  

    ↪label='Transpose Conv', linewidth=3)  

    if epochs_b:  

        axes[1, 0].plot(epochs_b, bilinear_history['val_loss'], 'r-',  

    ↪label='Bilinear Interp', linewidth=3)  

    axes[1, 0].set_title('Validation Loss Comparison', fontsize=14,  

    ↪fontweight='bold')  

    axes[1, 0].set_xlabel('Epoch')  

    axes[1, 0].set_ylabel('Validation Loss')  

    axes[1, 0].legend()  

    axes[1, 0].grid(True, alpha=0.3)

# Individual IoU comparison
if epochs:

```

```


        axes[1, 1].plot(epochs, transpose_history['val_iou'], 'b-', □
    ↵label='Transpose Conv', linewidth=3)
    if epochs_b:
        axes[1, 1].plot(epochs_b, bilinear_history['val_iou'], 'r-', □
    ↵label='Bilinear Interp', linewidth=3)
    axes[1, 1].set_title('Validation IoU Comparison', fontsize=14, □
    ↵fontweight='bold')
    axes[1, 1].set_xlabel('Epoch')
    axes[1, 1].set_ylabel('Mean IoU')
    axes[1, 1].legend()
    axes[1, 1].grid(True, alpha=0.3)

    # Individual accuracy comparison
    if epochs:
        axes[1, 2].plot(epochs, transpose_history['val_accuracy'], 'b-', □
    ↵label='Transpose Conv', linewidth=3)
        if epochs_b:
            axes[1, 2].plot(epochs_b, bilinear_history['val_accuracy'], 'r-', □
    ↵label='Bilinear Interp', linewidth=3)
        axes[1, 2].set_title('Validation Accuracy Comparison', fontsize=14, □
    ↵fontweight='bold')
        axes[1, 2].set_xlabel('Epoch')
        axes[1, 2].set_ylabel('Pixel Accuracy')
        axes[1, 2].legend()
        axes[1, 2].grid(True, alpha=0.3)

    plt.tight_layout()
    plt.show()

# Plot the training curves
plot_training_curves(transpose_history, bilinear_history)

```



1.6.2 2. Segmentation Results Visualization

```
[14]: def visualize_segmentation_results(model_transpose, model_bilinear, dataset, device, num_samples=3):
    """Visualize segmentation results from both models"""

    # Helper function to convert tensor to RGB image
    def tensor_to_rgb(tensor):
        # Denormalize
        mean = torch.tensor([0.485, 0.456, 0.406]).view(3, 1, 1)
        std = torch.tensor([0.229, 0.224, 0.225]).view(3, 1, 1)
        tensor = tensor * std + mean
        tensor = torch.clamp(tensor, 0, 1)
        return tensor.permute(1, 2, 0).numpy()

    def mask_to_rgb(mask):
        """Convert class mask to RGB using VOC colormap"""
        h, w = mask.shape
        rgb = np.zeros((h, w, 3), dtype=np.uint8)
        for class_id in range(len(VOC_COLORMAP)):
            rgb[mask == class_id] = VOC_COLORMAP[class_id]
        return rgb
```

```

model_transpose.eval()
model_bilinear.eval()

fig, axes = plt.subplots(num_samples, 5, figsize=(20, 4*num_samples))
if num_samples == 1:
    axes = axes.reshape(1, -1)

with torch.no_grad():
    for i in range(num_samples):
        # Get sample
        image, mask_gt = dataset[i]
        image_batch = image.unsqueeze(0).to(device)

        # Generate predictions
        output_transpose = model_transpose(image_batch)
        output_bilinear = model_bilinear(image_batch)

        # Get predicted classes
        _, pred_transpose = torch.max(output_transpose, 1)
        _, pred_bilinear = torch.max(output_bilinear, 1)

        # Convert to numpy
        image_np = tensor_to_rgb(image.cpu())
        mask_gt_np = mask_gt.cpu().numpy()
        pred_transpose_np = pred_transpose.squeeze().cpu().numpy()
        pred_bilinear_np = pred_bilinear.squeeze().cpu().numpy()

        # Convert masks to RGB
        mask_gt_rgb = mask_to_rgb(mask_gt_np)
        pred_transpose_rgb = mask_to_rgb(pred_transpose_np)
        pred_bilinear_rgb = mask_to_rgb(pred_bilinear_np)

        # Plot images
        axes[i, 0].imshow(image_np)
        axes[i, 0].set_title('Original Image', fontweight='bold')
        axes[i, 0].axis('off')

        axes[i, 1].imshow(mask_gt_rgb)
        axes[i, 1].set_title('Ground Truth', fontweight='bold')
        axes[i, 1].axis('off')

        axes[i, 2].imshow(pred_transpose_rgb)
        axes[i, 2].set_title('Transpose Conv', fontweight='bold')
        axes[i, 2].axis('off')

        axes[i, 3].imshow(pred_bilinear_rgb)
        axes[i, 3].set_title('Bilinear Interp', fontweight='bold')

```

```

        axes[i, 3].axis('off')

        # Difference visualization
        diff = np.abs(pred_transpose_np.astype(np.float32) - pred_bilinear_np.astype(np.float32))
        axes[i, 4].imshow(diff, cmap='hot')
        axes[i, 4].set_title('Prediction Difference', fontweight='bold')
        axes[i, 4].axis('off')

        # Calculate metrics for this sample
        iou_transpose = calculate_iou(torch.from_numpy(pred_transpose_np), torch.from_numpy(mask_gt_np))
        iou_bilinear = calculate_iou(torch.from_numpy(pred_bilinear_np), torch.from_numpy(mask_gt_np))

        acc_transpose = calculate_pixel_accuracy(torch.from_numpy(pred_transpose_np), torch.from_numpy(mask_gt_np))
        acc_bilinear = calculate_pixel_accuracy(torch.from_numpy(pred_bilinear_np), torch.from_numpy(mask_gt_np))

        print(f"Sample {i+1}:")
        print(f"  Transpose Conv - IoU: {iou_transpose:.4f}, Accuracy:{acc_transpose:.4f}")
        print(f"  Bilinear Interp - IoU: {iou_bilinear:.4f}, Accuracy:{acc_bilinear:.4f}")
        print()

        plt.tight_layout()
        plt.show()

# Visualize results on test samples
print("=" * 80)
print("SEGMENTATION RESULTS VISUALIZATION")
print("=" * 80)

visualize_segmentation_results(fcn_transpose, fcn_bilinear, test_dataset, device, num_samples=3)

```

```
=====
SEGMENTATION RESULTS VISUALIZATION
=====
```

Sample 1:

Transpose Conv - IoU: 0.6796, Accuracy: 0.8473
Bilinear Interp - IoU: 0.6842, Accuracy: 0.8323

Sample 2:


Transpose Conv - IoU: 0.7543, Accuracy: 0.8743

Bilinear Interp - IoU: 0.7921, Accuracy: 0.8937

Sample 3:

Transpose Conv - IoU: 0.7922, Accuracy: 0.8762

Bilinear Interp - IoU: 0.8193, Accuracy: 0.8806



1.6.3 3. Comparison Table

```
[15]: # Comprehensive comparison of both models
print("=" * 80)
print("COMPREHENSIVE MODEL COMPARISON")
print("=" * 80)

# Evaluate both models on the full test set
print("Evaluating models on test set...")

# Transpose convolution model
fcn_transpose.eval()
test_loss_transpose, test_iou_transpose, test_acc_transpose = evaluate_model(
    fcn_transpose, test_loader, nn.CrossEntropyLoss(ignore_index=255), device
)

# Bilinear interpolation model
fcn_bilinear.eval()
```

```

test_loss_bilinear, test_iou_bilinear, test_acc_bilinear = evaluate_model(
    fcn_bilinear, test_loader, nn.CrossEntropyLoss(ignore_index=255), device
)

# Calculate model parameters
def count_parameters(model):
    return sum(p.numel() for p in model.parameters() if p.requires_grad)

params_transpose = count_parameters(fcn_transpose)
params_bilinear = count_parameters(fcn_bilinear)

# Create comparison table
import pandas as pd

comparison_data = {
    'Metric': [
        'Test Loss',
        'Test Mean IoU',
        'Test Pixel Accuracy',
        'Model Parameters (M)',
        'Upsampling Method',
        'Memory Footprint',
        'Training Stability',
        'Inference Speed'
    ],
    'Transpose Convolution': [
        f'{test_loss_transpose:.4f}',
        f'{test_iou_transpose:.4f}',
        f'{test_acc_transpose:.4f}',
        f'{params_transpose/1e6:.2f}',
        'Learnable weights',
        'Higher (learnable params)',
        'More stable gradients',
        'Slightly slower'
    ],
    'Bilinear Interpolation': [
        f'{test_loss_bilinear:.4f}',
        f'{test_iou_bilinear:.4f}',
        f'{test_acc_bilinear:.4f}',
        f'{params_bilinear/1e6:.2f}',
        'Fixed interpolation',
        'Lower (no extra params)',
        'Standard',
        'Faster'
    ]
}

```

```

comparison_df = pd.DataFrame(comparison_data)

print("\nDETAILED COMPARISON TABLE")
print("=" * 80)
print(comparison_df.to_string(index=False))

# Additional analysis
print(f"\n\nPERFORMANCE ANALYSIS")
print("=" * 40)

print(f"Best performing model (IoU): ", end="")
if test_iou_transpose > test_iou_bilinear:
    print(f"Transpose Convolution ({test_iou_transpose:.4f} vs {test_iou_bilinear:.4f})")
    better_model = "Transpose Convolution"
    iou_improvement = ((test_iou_transpose - test_iou_bilinear) / test_iou_bilinear) * 100
else:
    print(f"Bilinear Interpolation ({test_iou_bilinear:.4f} vs {test_iou_transpose:.4f})")
    better_model = "Bilinear Interpolation"
    iou_improvement = ((test_iou_bilinear - test_iou_transpose) / test_iou_transpose) * 100

print(f"IoU improvement: {iou_improvement:.2f}%")

print(f"\nBest performing model (Accuracy): ", end="")
if test_acc_transpose > test_acc_bilinear:
    print(f"Transpose Convolution ({test_acc_transpose:.4f} vs {test_acc_bilinear:.4f})")
    acc_improvement = ((test_acc_transpose - test_acc_bilinear) / test_acc_bilinear) * 100
else:
    print(f"Bilinear Interpolation ({test_acc_bilinear:.4f} vs {test_acc_transpose:.4f})")
    acc_improvement = ((test_acc_bilinear - test_acc_transpose) / test_acc_transpose) * 100

print(f"Accuracy improvement: {acc_improvement:.2f}%")

print(f"\nParameter difference: {abs(params_transpose - params_bilinear):,} {parameters}")
print(f"Relative parameter increase: {((params_transpose - params_bilinear) / params_bilinear) * 100:.2f}%")

# Store results for summary

```

```

results_summary = {
    'transpose': {
        'loss': test_loss_transpose,
        'iou': test_iou_transpose,
        'accuracy': test_acc_transpose,
        'parameters': params_transpose
    },
    'bilinear': {
        'loss': test_loss_bilinear,
        'iou': test_iou_bilinear,
        'accuracy': test_acc_bilinear,
        'parameters': params_bilinear
    }
}

```

COMPREHENSIVE MODEL COMPARISON

Evaluating models on test set...

```

Evaluating: 100%|      | 100/100 [00:02<00:00, 47.86it/s, Loss=0.1351,
IoU=0.6449, Acc=0.9496]
Evaluating: 100%|      | 100/100 [00:02<00:00, 47.86it/s, Loss=0.1351,
IoU=0.6449, Acc=0.9496]
Evaluating: 100%|      | 100/100 [00:01<00:00, 50.29it/s, Loss=0.1064,
IoU=0.6782, Acc=0.9632]

```

DETAILED COMPARISON TABLE

Metric	Transpose Convolution	Bilinear Interpolation
Test Loss	0.4533	0.4271
Test Mean IoU	0.5819	0.6040
Test Pixel Accuracy	0.8782	0.8835
Model Parameters (M)	23.71	23.58
Upsampling Method	Learnable weights	Fixed interpolation
Memory Footprint	Higher (learnable params)	Lower (no extra params)
Training Stability	More stable gradients	Standard
Inference Speed	Slightly slower	Faster

PERFORMANCE ANALYSIS

Best performing model (IoU): Bilinear Interpolation (0.6040 vs 0.5819)
 IoU improvement: 3.80%

Best performing model (Accuracy): Bilinear Interpolation (0.8835 vs 0.8782)

Accuracy improvement: 0.60%

Parameter difference: 127,008 parameters

Relative parameter increase: 0.54%

1.7 Summary and Analysis

This task involved building, training, and evaluating a deep learning model called a Fully Convolutional Network (FCN) for semantic segmentation. The goal was to train the model to understand what object each pixel in an image belongs to (e.g., person, cat, train, or background).

The project used a pre-trained ResNet50 model as the “backbone” to extract image features. The key part of the experiment was to compare two different methods for upsampling (making the small feature map big again) to create the final pixel-level prediction:

1. **Transpose Convolution:** A learnable upsampling layer (`nn.ConvTranspose2d`) that learns the best way to enlarge the image.
2. **Bilinear Interpolation:** A fixed, non-learnable upsampling method (`nn.Upsample`) that uses a standard mathematical formula.

Both models were trained for 30 epochs on a subset of the Pascal VOC dataset (800 training images, 200 testing images) to compare their performance.

1.7.1 Analysis of Results

The two models were compared using training curves, final test metrics, and visual examples.

1. Training Curves

- **Training vs. Validation:** For both models, the training accuracy and IoU (Intersection over Union) were much higher than the validation metrics. This indicates that the models were overfitting to the training data, which is common.
- **Model Comparison:** The validation curves for the **Bilinear Interpolation** model were consistently better than the Transpose Convolution model. It achieved a slightly lower validation loss and a slightly higher validation IoU and accuracy throughout most of the training process.

2. Qualitative (Visual) Comparison Visual analysis of the test images shows that the **Bilinear Interpolation** model often produced cleaner and more accurate segmentation maps. For example, in the image of the train, the bilinear model correctly identified more of the train’s body, while the transpose convolution model missed a part of it. This suggests the simpler, non-learnable upsampling method produced better generalizations.

3. Quantitative Comparison The final test metrics, compiled in the table below, confirm the findings from the training curves and visual analysis.

Detailed Comparison Table

Metric	Transpose Convolution	Bilinear Interpolation
Test Mean IoU	0.5819	0.6040
Test Pixel Accuracy	0.8782	0.8835
Test Loss	0.4533	0.4271
Model Parameters (M)	23.71	23.58
Upsampling Method	Learnable weights	Fixed interpolation
Parameter Difference	+127,008 (0.54%)	(Baseline)

1.7.2 Conclusion

Based on all three types of analysis, the **Bilinear Interpolation** model performed better than the Transpose Convolution model.

- It achieved a **3.80% higher Mean IoU** and a **0.60% higher Pixel Accuracy** on the test set.
- It achieved this better performance while having **fewer parameters** (127,008 fewer) because its upsampling layers are not learnable.

This suggests that for this specific task and architecture, the simpler, fixed (non-learnable) bilinear upsampling method provided better results and was more efficient. The learnable transpose convolution layers may have added extra complexity, making the model slightly harder to train and more prone to overfitting, which led to slightly worse performance on the unseen test data.